


EHA Task Schematic

Claim, polluted evidence, structured output, and scoring layers



Core question: can the model avoid being captured by bad evidence, not merely answer correctly?